

Review of Computer Architecture

Glenn Bruns
CSUMB

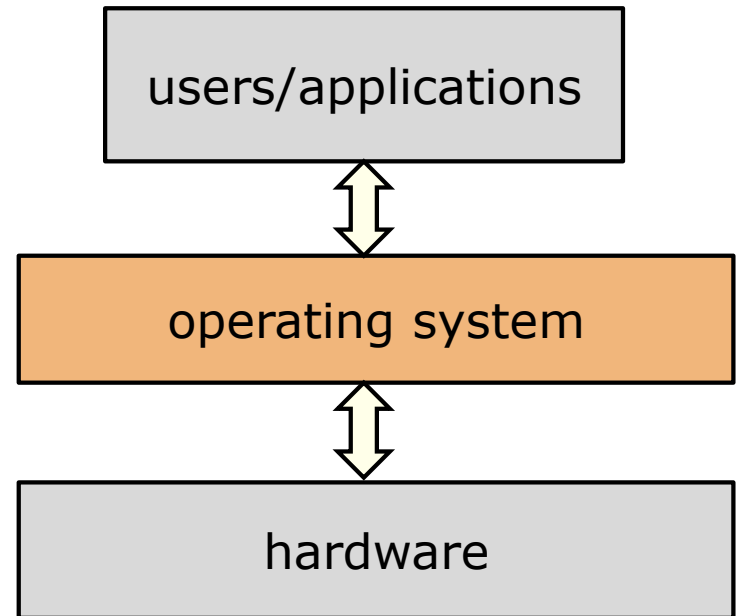
Lecture Objectives

At the end of this lecture, you should be able to:

- ☐ List the components of a computer
- ☐ Name some of the buses of a computer
- ☐ Describe a “memory hierarchy”
- ☐ Explain the layout of a program in memory

Motivation

- ❑ An OS is a software layer between the computer hardware and users/applications
- ❑ To understand what an OS does you need to know a little about the hardware

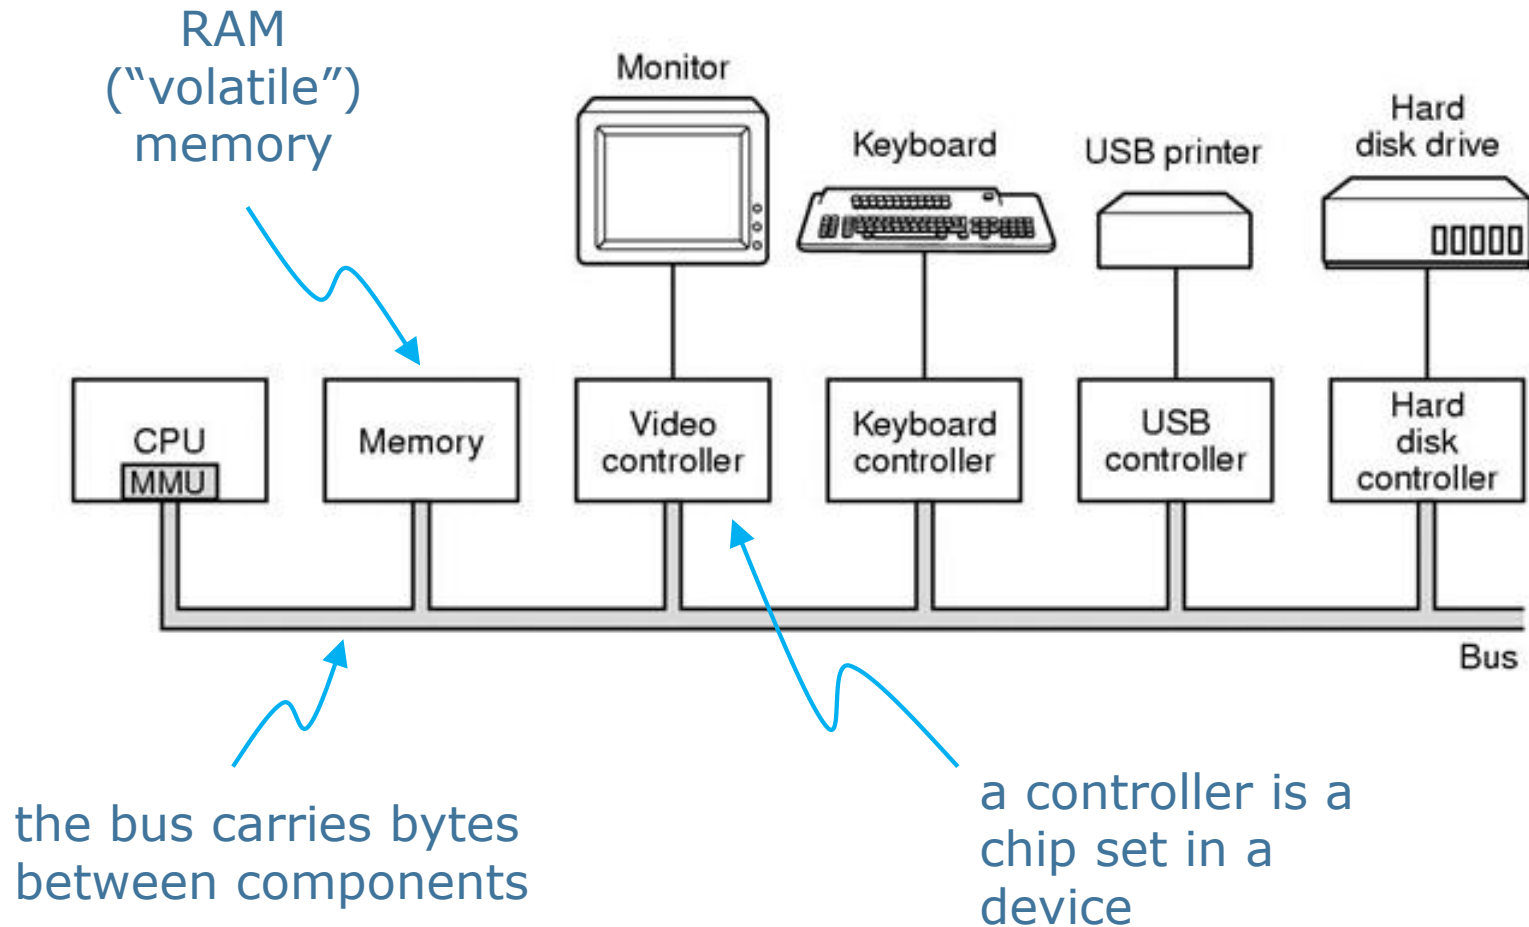


Components of a Computer

Question: what are the three main parts of a computer?

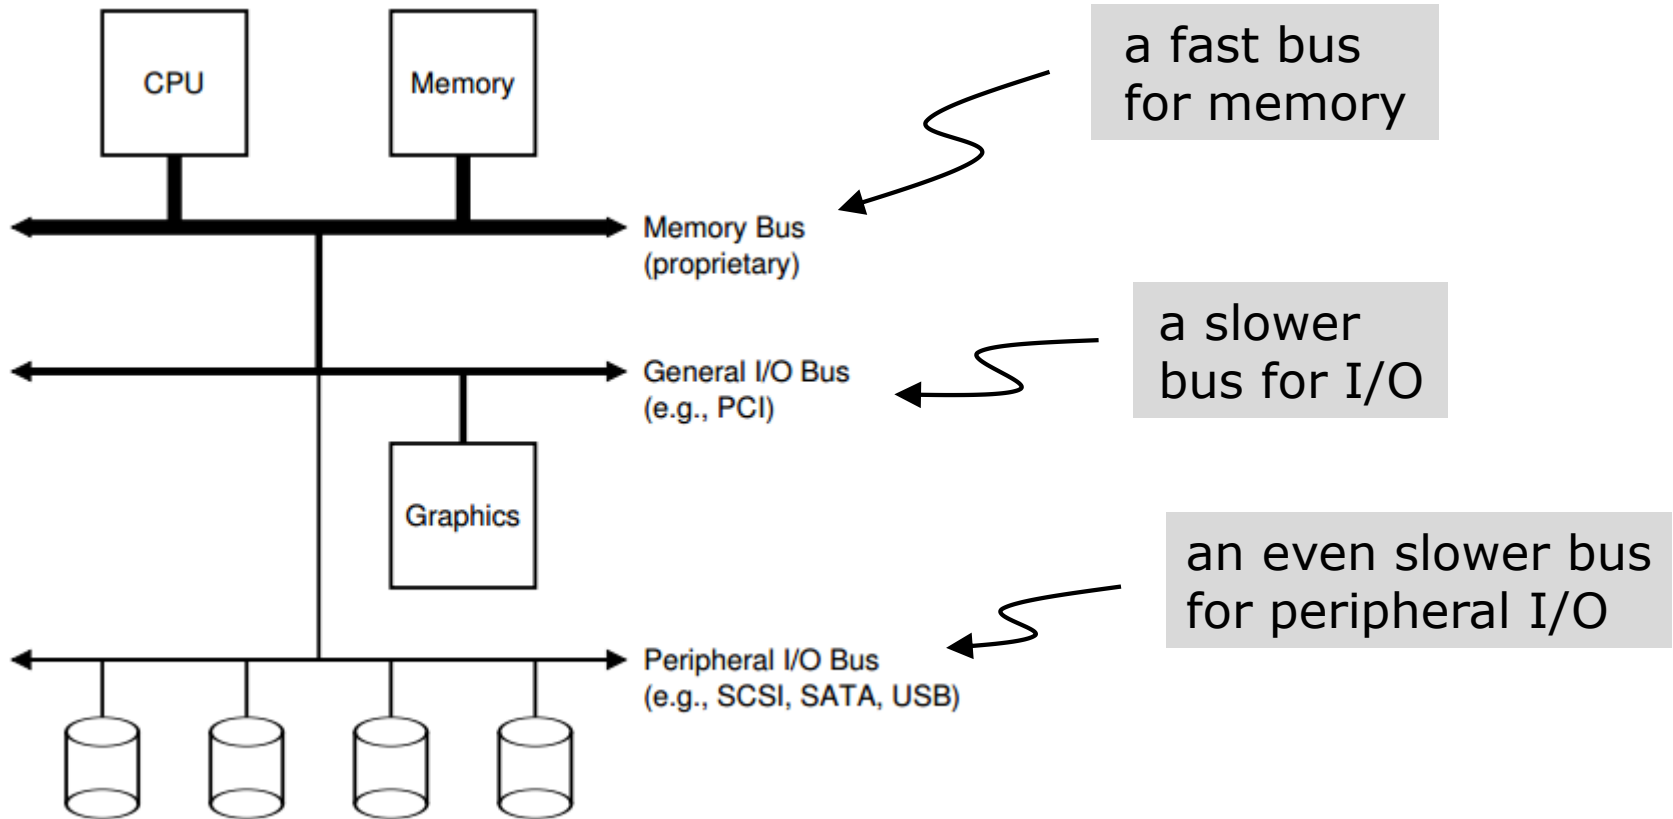
- ❑ Processor: performs computations
- ❑ Storage: stores data
- ❑ I/O devices: mouse, keyboard, printer, ...

Computer Architecture



(Figure from Tanenbaum, *Modern Operating Systems*, 3rd edition)

A little detail on buses



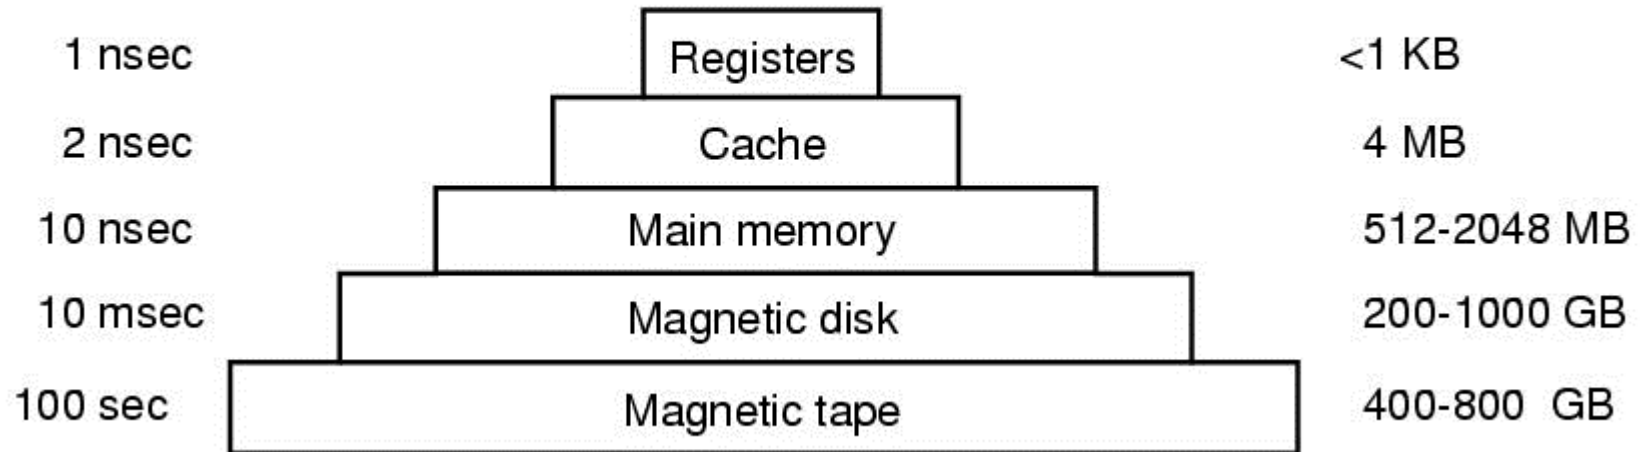
High performance components are closest to the CPU

(Figure from Operating Systems: Three Easy Pieces, Arpaci-Dusseau and Arpaci-Dusseau)

Computer storage hierarchy

Typical access time

Typical capacity



Higher layers are faster but have higher cost per bit.
Higher layers can be used to *cache* data in lower layers.

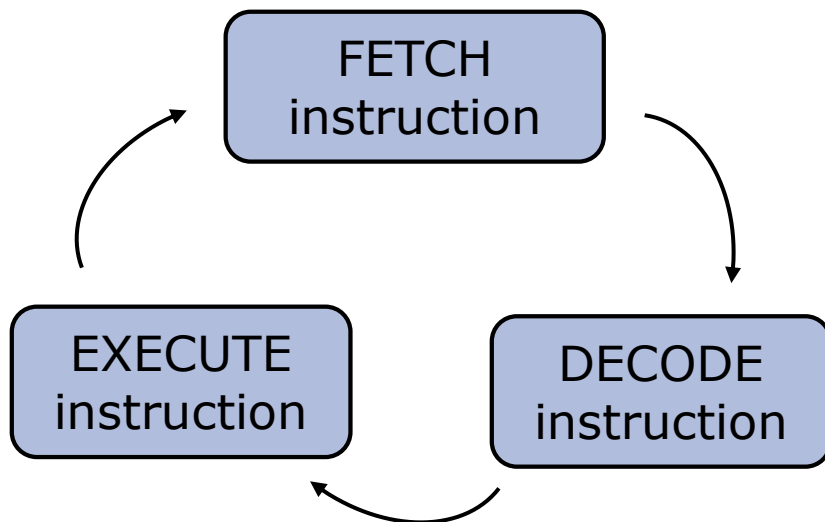
(Figure from Tanenbaum, *Modern Operating Systems*, 3rd edition)

Central Processing Unit (CPU)

Executes instructions in main memory

Contains a little memory (*registers*), including the *program counter* and *stack pointer*

CPU has a simple **instruction cycle**:



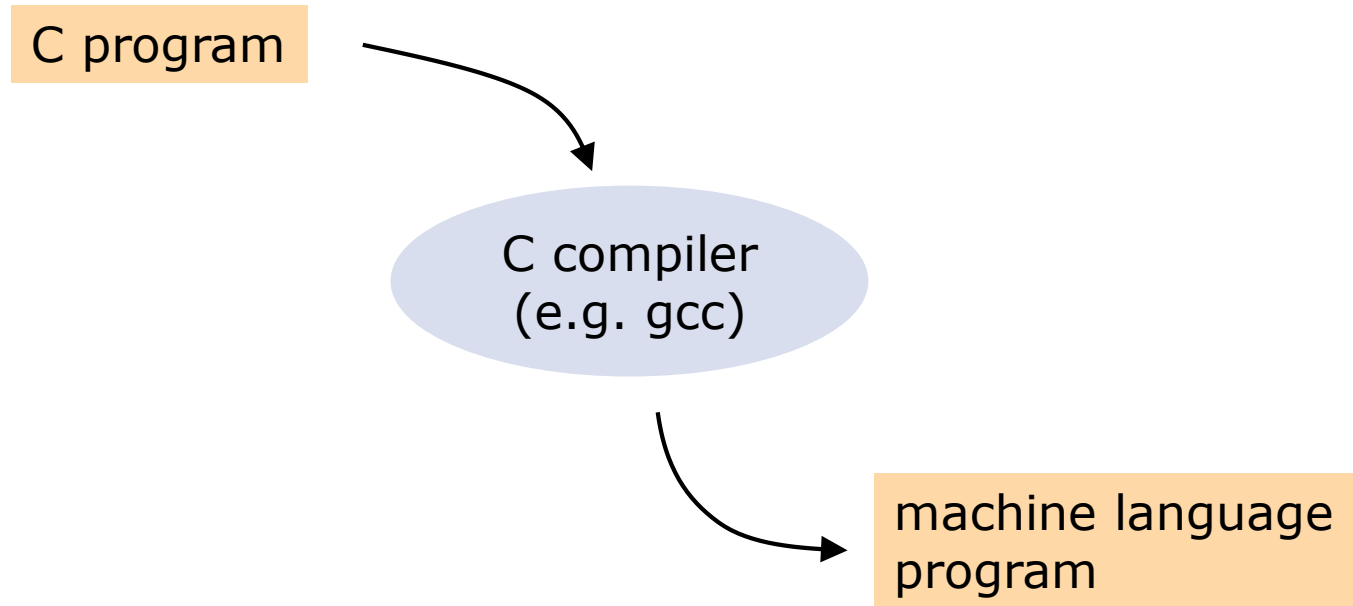
fetch: get instruction from memory location given by program counter

decode: get instruction arguments; break instruction down into parts

C Programming

- ❑ The language of Linux; also, the language of OS courses everywhere in the Milky Way galaxy.
- ❑ A simple and low-level language
- ❑ C is a subset of C++ (almost)
 - <http://stackoverflow.com/questions/1201593/c-subset-of-c-where-not-examples>

Machine language (aka binary)



- also called "binary programs", or a "program binary"
- super-low level programs – just a sequence of bytes
- the instructions are CPU specific
- **Instruction Set Architecture** (ISA) just means the set of instructions supported by a CPU

"Disassembling" a binary

```
int add_one_to(int iparam) {
    int ilocal = iparam + 1;
    return ilocal;
}
int main(int argc, char **argv) {
    int iresult = 0;
    iresult = add_one_to(iresult);
    return 0;
}
```



A tiny C program



Assembly code recreated
from binary for function
add_one_to()

```
gdb -q a.out
Reading symbols from /home/CLASSES/brunsglenn/return-oriented-prog/a.out...(no
debugging symbols found)...done.
```

```
(gdb) disas add_one_to
```

```
Dump of assembler code for function add_one_to:
```

```
0x08048394 <+0>:    push    %ebp
0x08048395 <+1>:    mov     %esp,%ebp
0x08048397 <+3>:    sub     $0x10,%esp
0x0804839a <+6>:    mov     0x8(%ebp),%eax
0x0804839d <+9>:    add     $0x1,%eax
0x080483a0 <+12>:   mov     %eax,-0x4(%ebp)
0x080483a3 <+15>:   mov     -0x4(%ebp),%eax
0x080483a6 <+18>:   leave
0x080483a7 <+19>:   ret
```

```
End of assembler dump.
```

ebp, esp, eax are
names of registers

'mov b,c' moves b to c

How to translate function calls?

- Where to store local variables of a function?
 - the memory used to store local variables isn't needed once the function returns
 - there can be many nested calls to a function
- Where to store the arguments of a function call?
- How does the function know where to return to?

```
// return true if a node in the
// tree has value key
int search_tree(TREE tree, int key) {
    int val; // local variable

    if (tree != NULL) {
        val = value(root(tree));
        if (val == key) {
            return true;
        }
        search_tree(left(tree), key);
        search_tree(right(tree), key);
    }
    return false;
}
```

Solution: use a stack

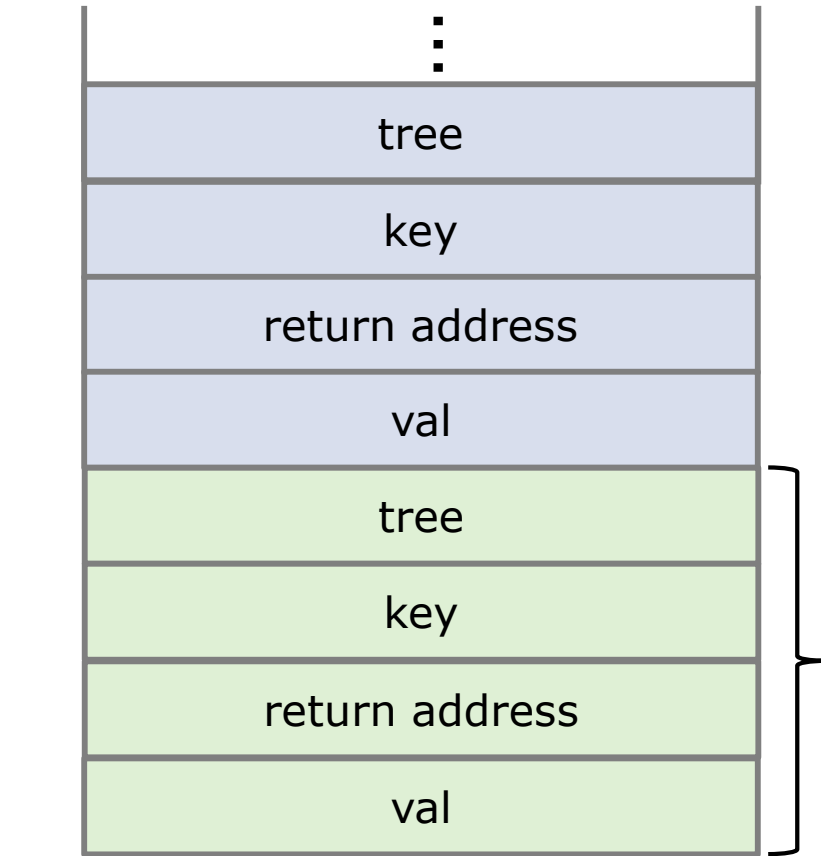
Suppose we have the call
`search_tree(t, 5)`

- ❑ the values of `t` and `key` are pushed onto a stack
- ❑ the address where `search_tree` should return to is pushed on the stack
- ❑ space is created on the stack for local variable `val`

```
// return true if a node in the
// tree has value key
int search_tree(TREE *tree, int key) {
    int val; // local variable

    if (tree != NULL) {
        val = value(root(tree));
        if (val == key) {
            return true;
        }
        search_tree(left(tree), key);
        search_tree(right(tree), key);
    }
    return false;
}
```

Solution: use a stack



```
// return true if a node in the
// tree has value key
int search_tree(TREE *tree, int key) {
    int val; // local variable

    if (tree != NULL) {
        val = value(root(tree));
        if (val == key) {
            return true;
        }
        search_tree(left(tree), key);
        search_tree(right(tree), key);
    }
    return false;
}
```

a stack frame

stack pointer

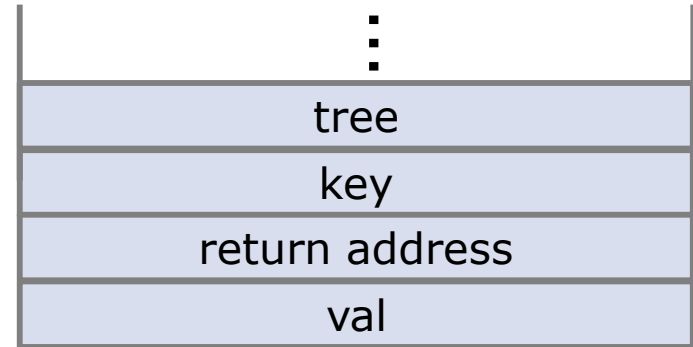
Using the stack

when `search_tree()` runs:

- ❑ it gets the value of `val` from the stack
- ❑ it gets the values of `tree` and `key` from the stack

when `search_tree()` returns:

- ❑ the stack frame is popped off the stack



```
// return true if a node in the
// tree has value key
int search_tree(TREE *tree, int key) {
    int val; // local variable

    if (tree != NULL) {
        val = value(root(tree));
        if (val == key) {
            return true;
        }
        search_tree(left(tree), key);
        search_tree(right(tree), key);
    }
    return false;
}
```

Looking at how the stack is used

```
int add_one_to(int iparam) {
    int ilocal = iparam + 1;
    return ilocal;
}
int main(int argc, char **argv) {
    int iresult = 0;
    iresult = add_one_to(iresult);
    return 0;
}
```

registers:

eax – general purpose

esp – pointer to top of stack

ebp – pointer to the stack frame

```
gdb -q a.out
```

```
Reading symbols from /home/CLASSES/brunsglenn/return-oriented-prog/a.out...(no
debugging symbols found)...done.
```

```
(gdb) disas add_one_to
```

```
Dump of assembler code for function add_one_to:
```

0x08048394 <+0>:	push	%ebp	# push ebp on stack to remember it
0x08048395 <+1>:	mov	%esp,%ebp	# point ebp to top of stack
0x08048397 <+3>:	sub	\$0x10,%esp	# make stack space for local vars
0x0804839a <+6>:	mov	0x8(%ebp),%eax	# copy iparam to eax register
0x0804839d <+9>:	add	\$0x1,%eax	# add 1 to the register
0x080483a0 <+12>:	mov	%eax,-0x4(%ebp)	# copy result to variable ilocal
0x080483a3 <+15>:	mov	-0x4(%ebp),%eax	# copy ilocal to eax register
0x080483a6 <+18>:	leave		# copy ebp to esp, restore ebp
0x080483a7 <+19>:	ret		# return control back to caller

```
End of assembler dump.
```


Memory layout for a running program

□ text segment

- stores the program's code
- also called 'code segment'

□ data segment

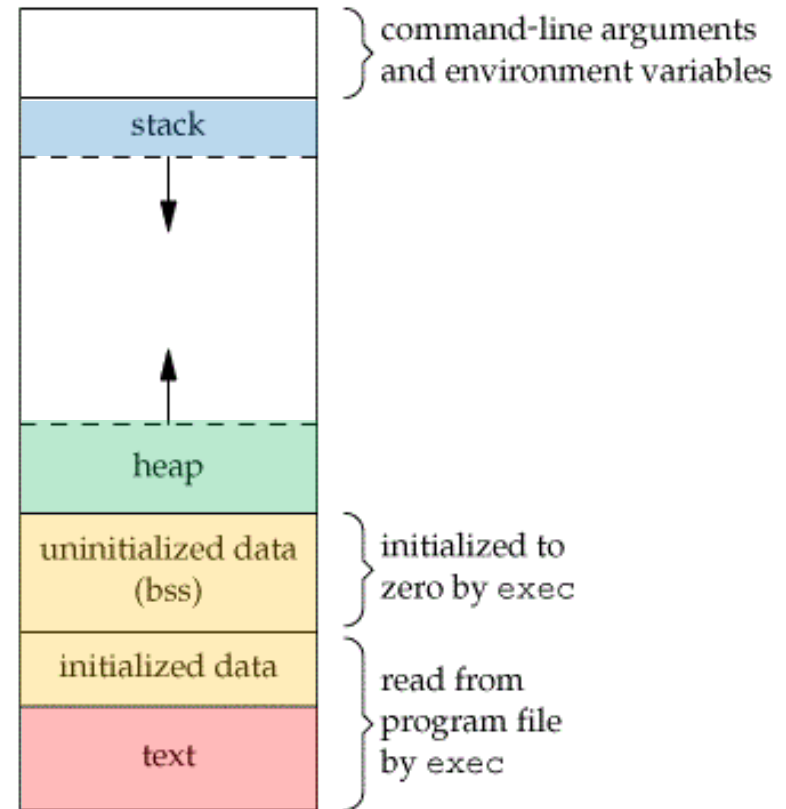
- used to store static variables (global vars, static local vars)

□ heap segment

- used to store dynamically-allocated variables

□ stack segment

- used to store program call stack
- "stack overflow" if too many nested calls



Summary

- ❑ The main components of a computer are the processor, memory, and I/O
- ❑ A computer has multiple buses, with different speeds
- ❑ Memory also comes in different speeds – faster memory is used to cache slower memory
- ❑ The four regions of a program in memory are: text segment, data segment, heap, stack