

Multi-level paging

Glenn Bruns
CSUMB

Lecture Objectives

After this lecture, you should be able to:

- ❑ Explain the reasons why multi-level paging is used
- ❑ Be able to manually perform address translation with multi-level paging

Recap: Problems with paging

- ❑ page tables are big
- ❑ page tables are slow



translation look-aside buffers
address the speed problem

Problem: big page tables

With 32 bit address space and 4 KB page size,
about 4 MB per page table


And one page table per process

A solution idea: make pages bigger

With a 32 bit address space, and a page size of 4 KB, a page table will have about a million entries.


How many entries if page size changed to 64 KB?

problem: “internal fragmentation”



Another solution Idea


- **insight:** processes usually use only a small part of their address space
- Most of a processes' page table is unused
- Break page table into **page-sized chunks**, use only some of these chunks



Page directory


A new **page directory** shows, for each group of pages, either:

- pointer to a page table “chunk”, or
- “not valid”



Good news: we only store the part of the page table we need
Bad news: page translation is more complicated

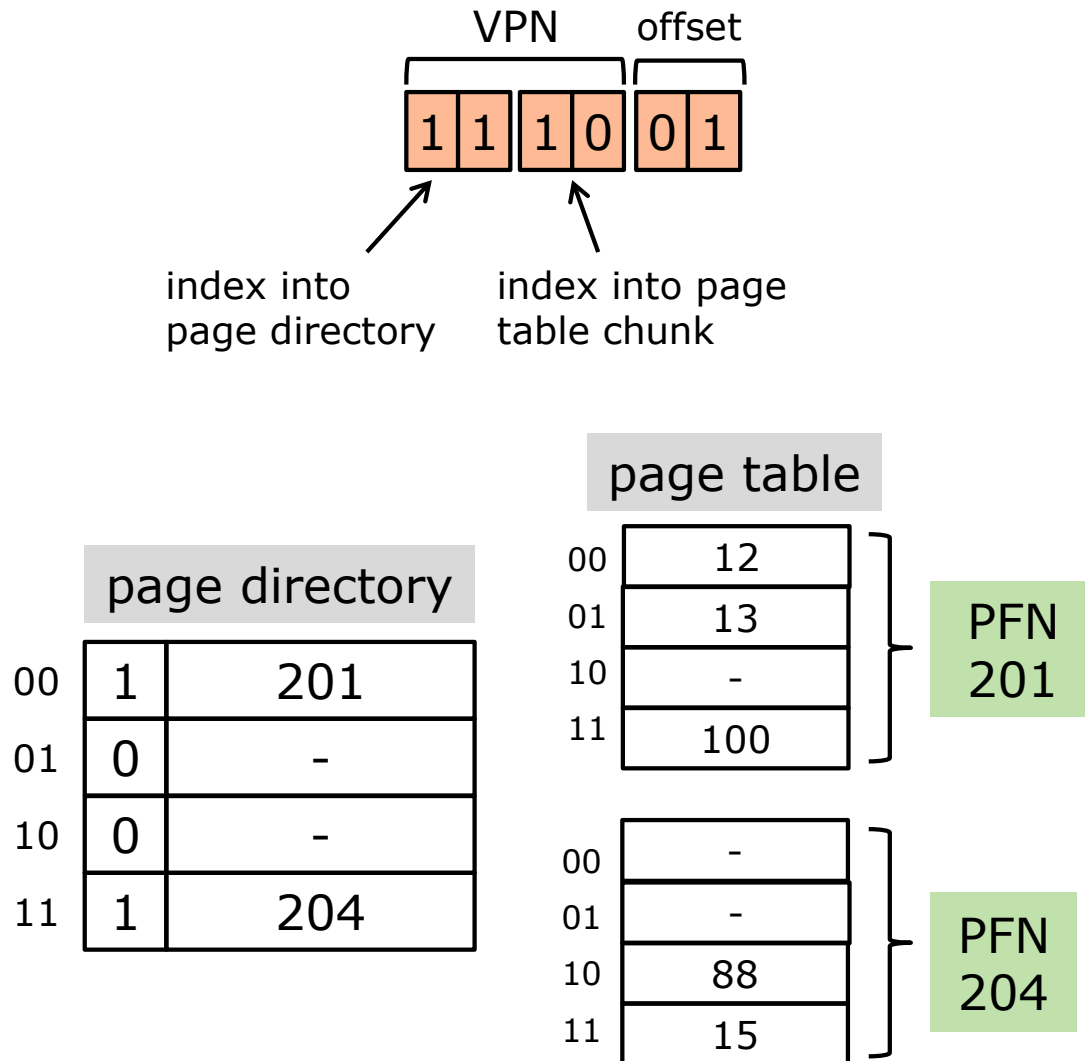
VPN now has 2 parts



page directory		
00	1	201
01	0	-
10	0	-
11	1	204

page table		
00	12	PFN 201
01	13	
10	-	
11	100	
00	-	PFN 204
01	-	
10	88	
11	15	

Address translation



1. Where is page directory?

Get page directory base address from a register

2. Which page table chunk?

Add page directory index to get page directory entry; get PFN of page table chunk

3. Which page?

If valid, add page table index to PFN of page table; get PFN of virtual page

Exercise

For each virtual address below, what is the physical address?

0 1 1 0 0 1

1 1 0 0 0 1

0 0 0 0 1 1

page directory


00	1	201
01	0	-
10	0	-
11	1	204

page table


00	12	PFN 201
01	13	
10	-	
11	100	
00	-	PFN 204
01	-	
10	88	
11	15	

How much space spacing?


Before: simple
page table



Look at page-
sized chunks



After: unused
chunks not stored




You save 1 page for every unused chunk, but add page directory

Address translation, including TLB

Use of TLB is the same as in the case of simple paging:

1. lookup the VPN part of the virtual address in the TLB
2. if a **TLB hit**:
 - form physical address from PFN in TLB entry
 - read value from physical address
3. if a **TLB miss**:
 - use the multi-level paging scheme to get a PFN value
 - stick the VPN, PFN pair in the TLB and retry instruction



see text for gory detail

Summary

- ❑ We need a page table for each process
- ❑ But most processes use only a small portion of their virtual address space
- ❑ Multi-level page tables avoid much of this waste