

Swapping

Glenn Bruns
CSUMB

Lecture Objectives

After this lecture, you should be able to:

- Explain the mechanics of swapping in an OS
- Manually perform some page replacement policies:
 - Belady's optimal policy, FIFO, LRU, Random
- Characterize the policies in terms of average memory time and cache hit rates

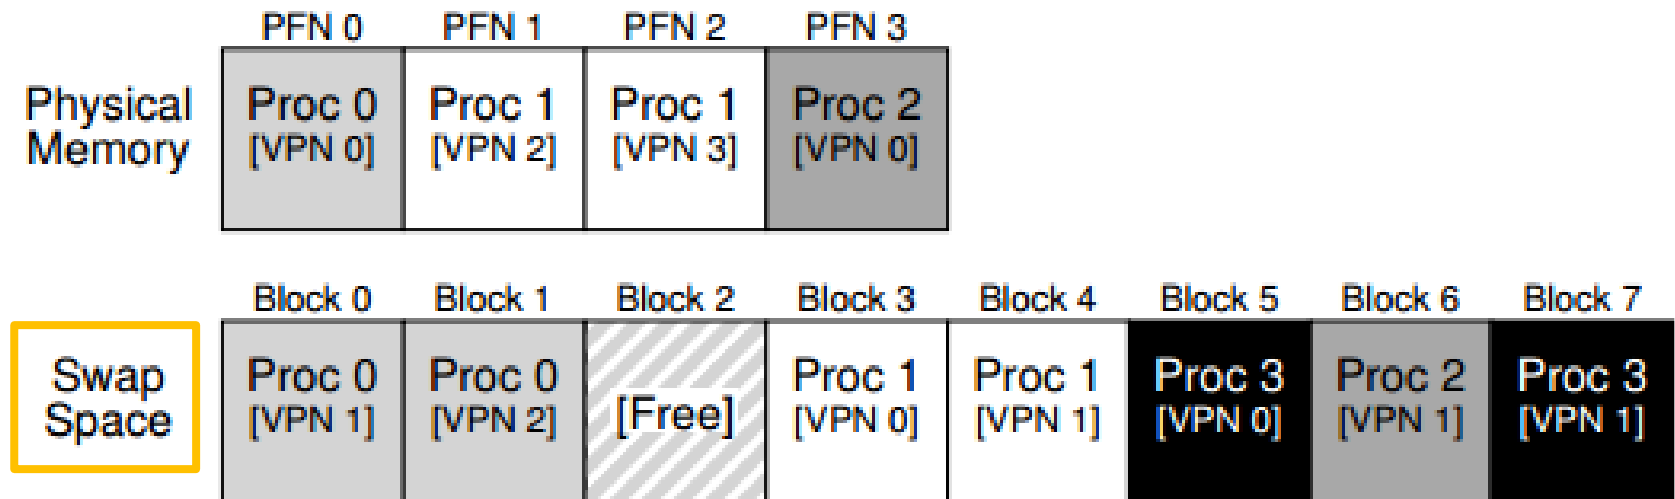
Problem: not enough memory

- ❑ OS may be running many processes at once
- ❑ Physical memory may not be large enough

Solution ideas

- ❑ Reduce size of address space available to processes
- ❑ Use disk – let programmers figure out how to move code and data from memory to disk
- ❑ Use disk – have OS figure out how to split address space between memory and disk

Pages in memory and on disk



In this tiny example, 4 pages of physical memory, and 8 pages of swap space.

Review: simple paging

Address translation process:

- use VPN as index into page table
- get physical frame number (PFN) from page table entry
- get physical address by combining PFN and offset part of virtual address

virtual address

<u>VPN</u>	<u>offset</u>
10	110010

page table

<u>VPN</u>	<u>PFN</u>
00	101
01	110
10	100
11	010

physical address

<u>PFN</u>	<u>offset</u>
100	110010

Review: simple paging with TLB

Address translation process:

Lookup PFN in TLB, using
VPN as key

if TLB miss:

- use VPN as index into page table
- read PFN from page table entry
- put PFN in TLB
- retry instruction

get physical address by
combining PFN and offset
part of virtual address

virtual address

<u>VPN</u>	<u>offset</u>
10	110010

TLB

(00, 101)
(10, 100)

page table

<u>VPN</u>	<u>PFN</u>
00	101
01	110
10	100
11	010

physical address

<u>PFN</u>	<u>offset</u>
100	110010

Mechanics of Swapping

- ❑ **present bit** in page table shows whether a page is present in memory
- ❑ hardware raises a **page fault** if a virtual page is not present

page fault handler:

- ❑ finds the disk address of the page in the page table
- ❑ issues a disk request to fetch the page into memory
- ❑ updates page table (once disk I/O completes)
- ❑ retries instruction

page table

00	1	12
01	0	120543
10	1	72
11	1	100

present bit

PFN **or** disk address

NEW!

Adding in swapping

Address translation process:

Lookup PFN in TLB, using VPN as key

if TLB miss:

 use VPN as index into page table

 if present bit of page table entry = 0:

 raise page fault

 else:

 read PFN from page table entry

 put PFN in TLB

 retry instruction

get physical address by combining PFN and offset part of virtual address

Swapping: policy

With swapping:

- we can think of memory as a cache for disk

When a page needs to be swapped in, but memory is full, we need to decide which page to swap out.

This is the **page replacement policy**.

This kind of policy is needed whenever caching is used.

FIFO policy

Evict the “oldest” cache element

	Access	Hit/Miss?	Evict	Resulting Cache State
	0			
	1			
	2			
3 element cache	0			
	1			
	3			
	0			
	3			
	1			
	2			
	1			

Least-recently used policy

Evict the least-recently used cache element

	Access	Hit/Miss?	Evict	Resulting Cache State
	0			
	1			
	2			
3 element cache	0			
	1			
	3			
	0			
	3			
	1			
	2			
	1			

Random policy

Evict a random cache element

3 element cache	Access	Hit/Miss?	Evict	Resulting Cache State
	0	Miss		0
	1	Miss		0, 1
	2	Miss		0, 1, 2
	0	Hit		0, 1, 2
	1	Hit		0, 1, 2
	3	Miss	0	1, 2, 3
	0	Miss	1	2, 3, 0
	3	Hit		2, 3, 0
	1	Miss	3	2, 0, 1
	2	Hit		2, 0, 1
	1	Hit		2, 0, 1

Optimal policy (Belady)

Evict cache element to be used furthest in future

	Access	Hit/Miss?	Evict	Resulting
				Cache State
3 element cache	0			
	1			
	2			
	0			
	1			
	3			
	0			
	3			
	1			
	2			
	1			

Why discuss optimal policy?

Optimal policy can't be used in realistic situations

But it serves as a good benchmark.

For example, if we do almost as well as optimal policy, we can't do much better.

Metrics for cache replacement policies

1. Cache hit rate
2. Average memory access time (AMAT)

Suppose we have 1000 memory accesses, and 950 result in cache hits.

What is the **Cache hit rate**?

Calculate the cache hit rates for the examples on past slides.

Average Memory Access Time

Suppose we have 1000 memory accesses and 950 of them are cache hits.

Also, it takes 100 ns to access memory (the cache), and 10 ms to access disk.

What's the **AMAT**?

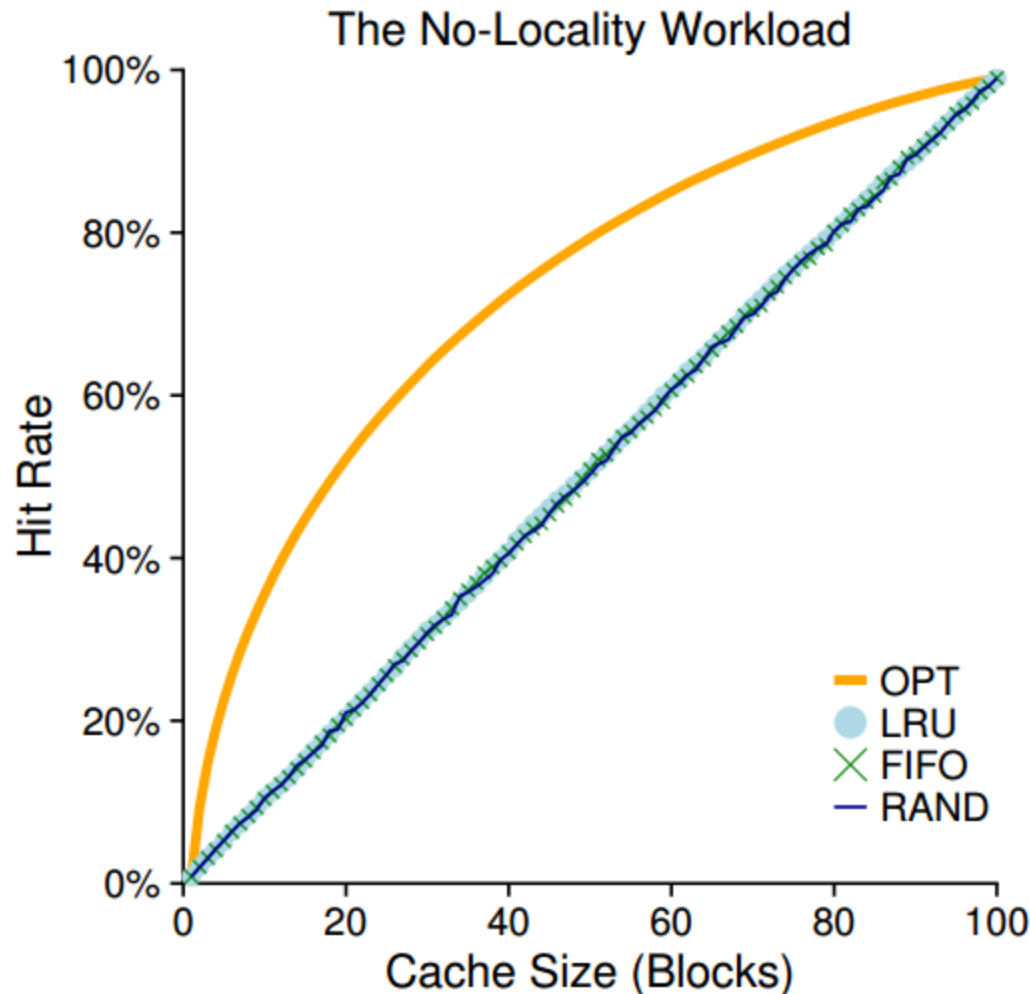
$$\text{AMAT} = (0.95 * 10 \text{ ns}) + (0.05 * 10 \text{ ms})$$

$$\sim 0.50 \text{ ms}$$

What if the cache hit rate is 99%?

$$\sim 0.10 \text{ ms}$$

Simulation with no-locality workload

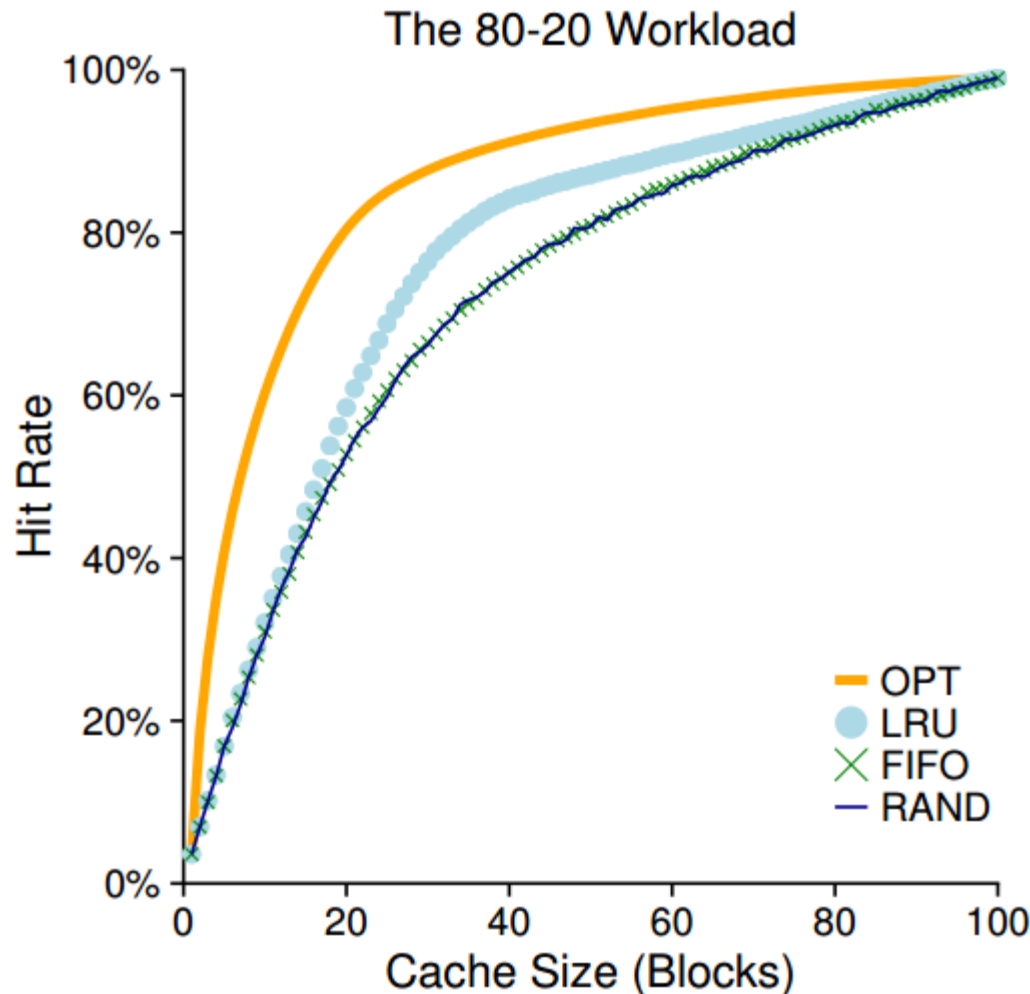


If memory accesses are completely random, all realistic policies have the same hit rate.

Cache size = 40 blocks → hit rate = 40%

(plot from OSTEP)

Simulation with 80:20 workload

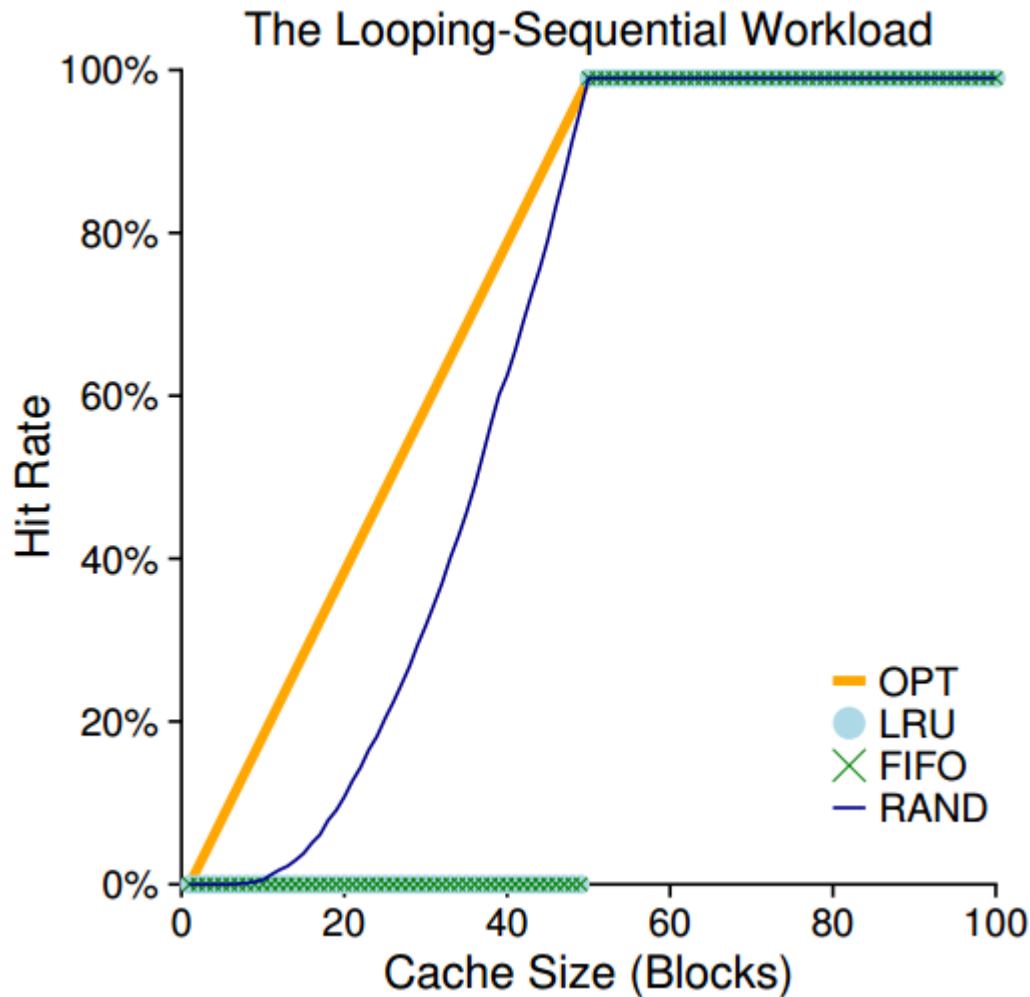


Some locality:
80% of
accesses made
to 20% of
memory.

At cache size
= 40 blocks,
LRU gives 82%
hit rate,
optimal gives
90%.

(plot from OSTEP)

Simulation with **looping workload**



Access pages 0,1,2,...,49 and repeat many times.

Shows strength of random policies.

(plot from OSTEP)

Summary

Swapping is used when insufficient memory for all current processes

Mechanism:

- a page fault is raised if the page table shows the page is not present in memory
- the page fault handler of the OS will then swap the page in from disk

Policy:

- In swapping (like in any situation with caching), we need a page replacement policy
- We looked at these page-replacement policies: FIFO, random, least-recently used (LRU), optimal
- Make sure you can do all these, plus LFU (see text)