

# *Concurrency and Threads*

---

Glenn Bruns  
CSUMB

# Lecture Objectives

---

After this lecture, you should be able to:

- ❑ Explain why concurrency is in this course
- ❑ Explain the benefits of concurrent programming
- ❑ Create a simple pthreads program

# Concurrency and Operating Systems

---

Concurrency issues were originally “discovered” in early OS work

Example:

- ☐ What if two processes want to write to the same file?
- ☐ How does the OS protect the file from getting corrupted?
- ☐ A big problem is that interrupts can happen at any time

# Benefits of concurrency

---

## Useful programming abstraction

- an IDE can handle editing with one thread and background compilation with another

## Responsiveness

- an app can handle user input while waiting for DB to initialize

## Leverage multicore machines and GPUs

- run separate code on separate cores

# Concurrent programming with processes

---

If you write a concurrent application using multiple processes, how do the processes interact?


**Threads** to the rescue!

Threads are lightweight processes that share virtual memory

# Threads vs. processes

---

- ❑ Threads and processes both have state
- ❑ Threads are faster to create, destroy, and context switch
- ❑ The threads within a process share a virtual address space
- ❑ But, each thread has its own stack, and its own registers



# C programming with threads

---

We'll use the C pthreads library

pthreads = "POSIX threads"

(POSIX = "portable OS interface")

Benefits:

- pthreads are widely used
- it's good to be familiar with primitives like locks and condition variables

Drawbacks:

- low level; easy to make mistakes

# Some multi-threaded C code

---

```
#include <stdio.h>
#include <assert.h>
#include <pthread.h> ←

void *mythread(void *arg) {
    printf("%s\n", (char *) arg);
    return NULL;
}

int
main(int argc, char *argv[]) {
    pthread_t p1, p2;
    int rc;
    printf("main: begin\n");
    rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
    rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
    // join waits for the threads to finish
    rc = pthread_join(p1, NULL); assert(rc == 0);
    rc = pthread_join(p2, NULL); assert(rc == 0);
    printf("main: end\n");
    return 0;
}
```



# Compiling and running

```
$ gcc -o t0 t0.c -pthread
```

```
$ ./t0
```

```
main: begin
```

```
A
```

```
B
```

```
main: end
```

```
$ ./t0
```

```
main: begin
```

```
A
```

```
B
```

```
main: end
```

```
$ ./t0
```

```
main: begin
```

```
B
```

```
A
```

```
main: end
```

```
$
```

use `-lpthread` if `-pthread`  
not available

non-determinism!

# pthread\_create()

---

create a new thread

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine)(void*),  
                  void *arg);
```

thread: pointer to a pthread\_t structure  
attr: thread attributes (or NULL, for defaults)  
start\_routine: function to start running  
arg: arguments for start\_routine (NULL if no args)

start\_routine has input and output of type void \*

```
void *my_thread(void * args) { ... }
```


# Thread creation example

```
typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *)arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int
main(int argc, char *argv[]) {
    pthread_t p;
    int rc;
    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

cast the argument to  
the right type



# pthread\_join()

---

sleep until the given thread dies (a blocking call)

```
int pthread_join(pthread_t thread,  
                 void **value_ptr);
```

thread: a pthread\_t structure (not a pointer to one!)

value\_ptr: pointer to the expected return value

# Threads that share data

```
static volatile int counter = 0;
```

shared variable

```
void *mythread(void *arg) {  
    printf("%s: begin\n", (char *) arg);  
    int i;  
    for (i = 0; i < 1e7; i++) {  
        counter = counter + 1;  
    }  
    printf("%s: done\n", (char *) arg);  
    return NULL;  
}
```

each thread runs mythread()

```
int main(int argc, char *argv[]) {  
    pthread_t p1, p2;  
    printf("main: begin (counter = %d)\n", counter);  
    Pthread_create(&p1, NULL, mythread, "A");  
    Pthread_create(&p2, NULL, mythread, "B");  
  
    // join waits for the threads to finish  
    Pthread_join(p1, NULL);  
    Pthread_join(p2, NULL);  
    printf("main: done with both (counter = %d)\n", counter);  
    return 0;  
}
```

Question:  
what exactly  
will this  
program  
output?

# Compiling and running


---

```
$ gcc -o t1 t1.c -pthread
$ ./t1
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 13505135)
$
$ ./t1
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 11928697)
$
```

# Interleaved program execution

---

```
for (i = 0; i < 1e7; i++) {  
    counter = counter + 1;  
}
```



mov	0x8049d50,%eax	# copy counter value into eax register
add	\$0x1,%eax	# increment value in eax register
mov	%eax,0x8049d50	# copy eax register value to counter

## **thread 1**

load counter into eax  
increment eax

## **thread 2**

load counter into eax  
increment eax  
store eax as counter

store eax as counter

*context switch*

*context switch*

# Solution

---

- ❑ Use “synchronization primitives” to control when threads can run
- ❑ But OS process scheduling also wants to control the execution of threads!



# Summary

---

- ❑ We're now in the realm of multi-threaded programs
- ❑ Pros:
  - can leverage multi-core
  - many programs are naturally concurrent
- ❑ Cons:
  - concurrency bugs, non-determinism
- ❑ We'll learn to write multi-threaded C, and to use pthread synchronization primitives