

# *Lock-based Concurrent Data Structures*

---

Glenn Bruns  
CSUMB

# Lecture Objectives

---

After this lecture, you should be able to:

- ❑ Build thread-safe data structures using locks
- ❑ Explain performance issues when locks are added to data structures

# Recall: race conditions

---

A **critical section** is a piece of code that accesses a shared resource, such as a shared variable.

When multiple threads are running, and the output depends on the timing of their execution, then the code has a **race condition**.

Avoid race conditions by allowing one thread at a time in the critical section.

# Counter code with pthread locks

---

```
static volatile int counter = 0;
pthread_mutex_t lock;

void *mythread(void *arg) {
    int i;
    for (i = 0; i < 1e7; i++) {
        Pthread_mutex_lock(&lock);
        counter = counter + 1;
        Pthread_mutex_unlock(&lock);
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    Pthread_mutex_init(&lock, NULL);

    pthread_t p1, p2;
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");

    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("main: counter = %d\n", counter);
    return 0;
}
```

This is ugly.

The “application” code is aware of the lock, and has to use it correctly.

It would be better to have a “thread-safe” counter.

# A thread-safe counter object

---

```
typedef struct __counter_t {
    int value;
    pthread_mutex_t lock;
} counter_t;

void init(counter_t *c) {
    c->value = 0;
    Pthread_mutex_init(&c->lock, NULL);
}

void increment(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value++;
    Pthread_mutex_unlock(&c->lock);
}

void decrement(counter_t *c) {
    // similar to increment
}

int get(counter_t *c) {
    // similar to increment
    // a value must be returned
}
```

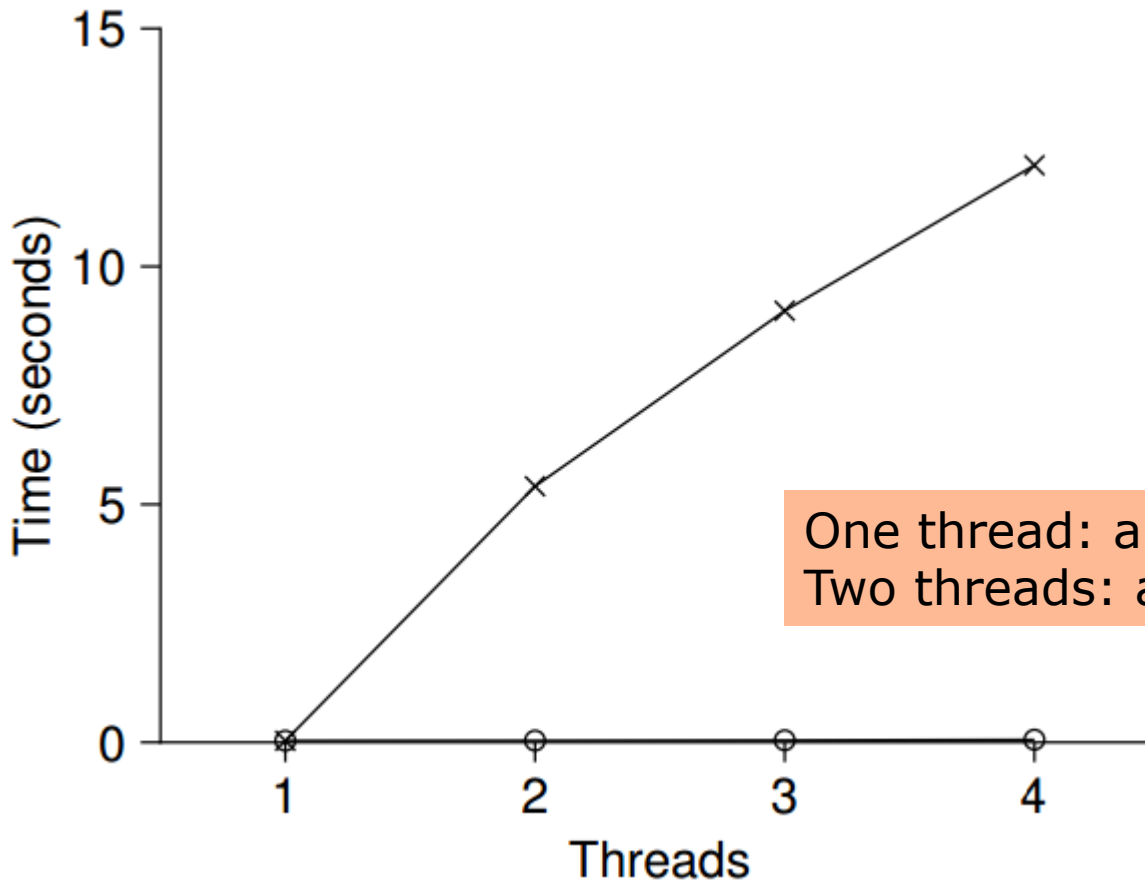
Now we have a counter object that includes locks.

Methods are **init**, **increment**, **decrement**, and **get**.

Users aren't aware of the locks, but the code works correctly with threads.  
(it's **thread-safe**)

# Counter performance

---



One thread: about 0.03 seconds  
Two threads: about 5 seconds

(plot from Operating Systems: Three Easy Pieces, Arpaci-Dusseau et al)

# A concurrent linked list (part 1)

---

```
// a list node
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

// a list (one used per list)
typedef struct __list_t {
    node_t *head;
    pthread_mutex_t lock;
} list_t;

void List_Init(list_t *L) {
    L->head = NULL;
    pthread_mutex_init(&L->lock, NULL);
}
```

a node consists of a value and a pointer to the next node

a list consists of a lock and a pointer to the first node

# A concurrent linked list (part 2)

---

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; // fail
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; // success
}
```

Idea: lock the entire insert operation



# A concurrent linked list (part 3)

---

```
int List_Lookup(list_t *L, int key) {  
    pthread_mutex_lock(&L->lock);  
    node_t *curr = L->head;  
    while (curr) {  
        if (curr->key == key) {  
            pthread_mutex_unlock(&L->lock);  
            return 0; // success  
        }  
        curr = curr->next;  
    }  
    pthread_mutex_unlock(&L->lock);  
    return -1; // failure  
}
```

Idea: similar to  
List\_Insert

# Scaling list performance

---

- ❑ The insert and lookup operations use “coarse-grained” locking.
- ❑ Each thread that wants access gets a lock on the **entire list**.
- ❑ Would it help to lock each node individually?
  - threads could operate on different parts of the list at the same time
  - but a lot more locking/unlocking overhead
- ❑ Start with simple implementation; refine if necessary.

# Summary

---

- ❑ It's better to have locks hidden in data structures than handled by applications.
- ❑ Locks can be added to data structures to make them “thread-safe”.
- ❑ Often it's simple to add locks.
- ❑ In designing a thread-free data structure, remember that more concurrency may not mean better performance.