

# *Concurrency review*

---

Glenn Bruns  
CSUMB

# Exam structure

---

50 minutes

Concurrency concepts

Synchronization primitives

Pthreads programming

Command-line (bash, awk, etc.)

Process mgmt. and memory mgmt.

about 60%

about 20%

about 20%

You should be able to write and understand pthreads code.

# Concurrency concepts

---

concurrency

process

thread

shared memory

non-determinism

deadlock

critical section

race condition

mutual exclusion

lock

blocking operation

thread-safe data structure

condition variable

spurious wakeup

semaphore

synchronization primitive

synchronization variable

shared data

bounded buffer

synchronization

serialization


execution path

atomicity

# Threads and processes

- ❑ A process contains one or more threads
- ❑ Threads and processes both have state
- ❑ Threads are faster to create, destroy, and context switch
- ❑ The threads within a process share a virtual address space
- ❑ But, each thread has its own stack, and its own registers

a “classic”, single-threaded process



# Thread creation

---

```
int  
pthread_create(pthread_t *thread,  
               const pthread_attr_t *attr,  
               void *(*start_routine)(void*),  
               void *arg);
```

thread: pointer to a pthread\_t structure

attr: thread attributes (or NULL, for defaults)

start\_routine: function to start running

arg: arguments for start\_routine  
(NULL if no args)

# Thread completion

---

```
int  
pthread_join(pthread_t thread,  
              void **value_ptr);
```

thread:            a pthread\_t structure (not a pointer to one!)

value\_ptr:        pointer to the expected return value

# Race conditions and mutual exclusion

---

A **critical section** is a piece of code that accesses a shared resource, such as a shared variable.

When multiple threads are running, and the output depends on the timing of their execution, then the code has a **race condition**.

What we want is that at most one thread at a time can be in the critical section – this is **mutual exclusion**.

# Question

---

Two processes may share memory if they each have multiple threads

- a) true
- b) false

This is generally false, but there is an exception: processes can share memory if parts of their virtual memory spaces are mapped to the same area of physical memory.

We said this can happen with segmented memory, when two processes sharing running the same program can have their code regions mapped to the same physical memory.



# Question

---

Processes can have zero threads

- a) true
- b) false

False. Every process has at least one thread.

# Lock operations

---

- ❑ A lock can be in one of two states: busy or free
- ❑ A lock is initially in the free state
- ❑ Operation **lock** waits until the lock is free, and automatically make the lock busy
- ❑ Operation **unlock** makes the lock free.

'lock' is sometimes called 'acquire'  
'unlock' is sometimes called 'free'

'lock' is a blocking operation

# Using locks for mutual exclusion

---

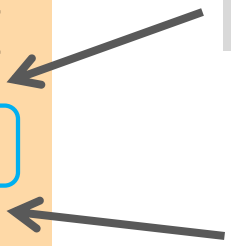
A **lock** is an object that can be “held” by only one thread at a time.

A **lock** can be acquired by a thread.

A **lock** can be released by a thread.

```
void *mythread(void *arg) {  
    int i;  
    for (i = 0; i < 1e7; i++) {  
        counter = counter + 1;  
    }  
    return NULL;  
}
```

acquire lock here

Two arrows point from the text boxes to the code. The first arrow points to the opening curly brace of the for loop, and the second arrow points to the closing curly brace of the for loop.

release lock here

critical section

An arrow points from the text 'critical section' to the line 'counter = counter + 1;' which is enclosed in a blue box.

# Question

What would happen if we acquired and released the lock at these locations? (assume two threads)

```
void *mythread(void *arg) {  
    int i;  
    for (i = 0; i < 1e7; i++) {  
        counter = counter + 1;  
    }  
  
    return NULL;  
}
```

acquire lock here

release lock here

critical section

One thread would do all of its counter increments, then the other thread would do all of its counter increments.

# Pthreads locks

```
pthread_mutex_t lock;
```



declare lock

```
void *mythread(void *arg) {  
    int i;  
    for (i = 0; i < 1e7; i++) {  
        Pthread_mutex_lock(&lock);  
        counter = counter + 1;  
        Pthread_mutex_unlock(&lock);  
    }  
    return NULL;  
}
```



acquire lock



release lock

```
int main(int argc, char *argv[]) {  
    Pthread_mutex_init(&lock);  
    pthread_t p1, p2;  
    Pthread_create(&p1, NULL, mythread, "A");  
    Pthread_create(&p2, NULL, mythread, "B");
```



initialize lock

```
    // join waits for the threads to finish  
    Pthread_join(p1, NULL);  
    Pthread_join(p2, NULL);  
    printf("main: done with both (counter = %d)\n", counter);  
    return 0;  
}
```

# Counter code with pthread locks

---

```
static volatile int counter = 0;
pthread_mutex_t lock;

void *mythread(void *arg) {
    int i;
    for (i = 0; i < 1e7; i++) {
        Pthread_mutex_lock(&lock);
        counter = counter + 1;
        Pthread_mutex_unlock(&lock);
    }
    return NULL;
}

int main(int argc, char *argv[]) {
    Pthread_mutex_init(&lock, NULL);

    pthread_t p1, p2;
    Pthread_create(&p1, NULL, mythread, "A");
    Pthread_create(&p2, NULL, mythread, "B");

    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("main: counter = %d\n", counter);
    return 0;
}
```

This is ugly.

The “application” code is aware of the lock, and has to use it correctly.

It would be better to have a “thread-safe” counter.

# A counter with locks

---

```
typedef struct __counter_t {
    int value;
    pthread_mutex_t lock;
} counter_t;

void init(counter_t *c) {
    c->value = 0;
    Pthread_mutex_init(&c->lock, NULL);
}

void increment(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value++;
    Pthread_mutex_unlock(&c->lock);
}

void decrement(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    c->value--;
    Pthread_mutex_unlock(&c->lock);
}

int get(counter_t *c) {
    Pthread_mutex_lock(&c->lock);
    int rc = c->value;
    Pthread_mutex_unlock(&c->lock);
    return rc;
}
```


Now we have a counter object that includes locks.

Operations are init, increment, decrement, and get.

Users aren't aware of the locks, but it works correctly with threads.

# Counter performance

---



One thread: about 0.03 seconds  
Two threads: about 5 seconds

(plot from Operating Systems: Three Easy Pieces, Arpaci-Dusseau et al)



# Question

---

All the threads of a process run the same code

- a) true
- b) false

False. Take for example the readers/writers problem. The readers and writers are threads in the same process.

# Question

---

The function 'pthread\_create' creates a clone of the currently running thread (like 'fork' of the process API).

- a) true
- b) false

False. One of the parameters of pthread\_create specifies the function that should be run by the newly-created thread.

# Condition variables

---

A **condition variable** is a synchronization object that lets threads wait efficiently

**pthread\_cond\_wait(cond, lock)**

- lock is released, calling thread is suspended and put on the condition variable's waiting list
- lock is re-acquired before wait returns

**pthread\_cond\_signal(cond)**

- takes a thread off the waiting list and marks it as "ready"
- if no thread on the waiting list, does nothing

**pthread\_cond\_broadcast(cond)**

- like signal, but takes **all** threads off the waiting list

# Question

---

When a lock is initialized, it is in the 'free' state.

- a) true
- b) false

True.

# Question

---

Condition variables can be implemented with locks

- a) true
- b) false

False. The only state of a lock is whether it is busy or free. Condition variables have a queue to keep track of waiting threads.

# Question

---

Which condition variable operations are blocking?

Only wait. A thread will never wait on signal.

# Classic concurrency problems


---

- ❑ Bounded buffer
- ❑ Readers/writers lock
- ❑ Synchronization barrier

# The bounded buffer problem

---

- ❑ One thread waits for buffer to be non-full before writing
- ❑ Another thread waits for buffer to be non-empty before reading
- ❑ Linux pipes use a bounded buffer





# Dahlin Method, part 1: class design

---

- ❑ Identify objects
- ❑ Define interfaces, and identify state variables
- ❑ Implement methods

This is just how you would do class design with an OO language

## Part 2: multi-threaded case

---

- ❑ add a **lock**
- ❑ add code to acquire and release the lock
- ❑ add zero or more **condition variables**
- ❑ add wait calls within loops
- ❑ add signal and broadcast calls

# Simple bounded buffer: class design

---

```
typedef struct {  
    // state variables  
    int cnt;        // 0 or 1, depending on whether buffer empty or not  
    int val;        // value of item in buffer  
} SBUF;  
  
// create a new synchronized buffer  
SBUF *sbuf_create();  
  
// write to the buffer  
void sbuf_write(SBUF *sbuf, int a);  
  
// read from the buffer  
int sbuf_read(SBUF *sbuf);
```

This illustrates how to do OO-style code (without inheritance) in C

# Implement methods

---

```
typedef struct {
    int cnt;
    int val;
} SBUF;

// create a new synchronized buffer
SBUF *sbuf_create() {
    SBUF *sbuf = (SBUF *)malloc(sizeof(SBUF));
    sbuf->cnt = 0;
    sbuf->val = 0;
}

// write to the buffer
void sbuf_write(SBUF *sbuf, int a) {
    sbuf->val = a;
    sbuf->cnt = 1;
}

// read from the buffer
int sbuf_read(SBUF *sbuf) {
    int a = sbuf->val;
    sbuf->cnt = 0;
    return(a);
}
```

# Add lock, and code to use it

---

```
typedef struct {  
    // state variables  
    int cnt;  
    int val;  
  
    // sync. variables  
    pthread_mutex_t lock;  
} SBUF;
```

```
void sbuf_write(SBUF *sbuf, int a) {  
    pthread_mutex_lock(&sbuf->lock);  
    sbuf->val = a;  
    sbuf->cnt = 1;  
    pthread_mutex_unlock(&sbuf->lock);  
}  
  
int sbuf_read(SBUF *sbuf) {  
    pthread_mutex_lock(&sbuf->lock);  
    int a = sbuf->val;  
    sbuf->cnt = 0;  
    pthread_mutex_unlock(&sbuf->lock);  
    return(a);  
}
```

- normally a shared object will have exactly one lock
- each method begins and end with locking/unlocking

# Add condition variables

---

```
typedef struct {  
    // state variables  
    int cnt;  
    int val;  
    // sync. variables  
    pthread_mutex_t lock;  
    pthread_cond_t read_go;  
    pthread_cond_t write_go;  
} SBUF;
```

- think about situations in which methods will have to wait
- if method never need to wait, no condition vars needed
- in this step the designer has a lot of freedom

# Add wait calls inside loops

```
void sbuf_write(SBUF *sbuf, int a) {
    pthread_mutex_lock(&sbuf->lock);

    while ("buffer full") {
        pthread_cond_wait(&sbuf->write_go, &sbuf->lock);
    }
    sbuf->val = a;
    sbuf->cnt = 1;

    pthread_mutex_unlock(&sbuf->lock);
}
```

Why are wait calls within loops?

```
int sbuf_read(SBUF *sbuf) {
    pthread_mutex_lock(&sbuf->lock);

    while ("buffer empty") {
        pthread_cond_wait(&sbuf->read_go, &sbuf->lock);
    }
    int a = sbuf->val;
    sbuf->cnt = 0;

    pthread_mutex_unlock(&sbuf->lock);
    return(a);
}
```

# Add signal and broadcast calls

```
void sbuf_write(SBUF *sbuf, int a) {
    pthread_mutex_lock(&sbuf->lock);

    while ("buffer full") {
        pthread_cond_wait(&sbuf->write_go, &sbuf->lock);
    }
    sbuf->val = a;
    sbuf->cnt = 1;

    pthread_mutex_unlock(&sbuf->lock);
}
```

- what happens if a thread signals "too much"

```
int sbuf_read(SBUF *sbuf) {
    pthread_mutex_lock(&sbuf->lock);

    while ("buffer empty") {
        pthread_cond_wait(&sbuf->read_go, &sbuf->lock);
    }
    int a = sbuf->val;
    sbuf->cnt = 0;

    pthread_mutex_unlock(&sbuf->lock);
    return(a);
}
```

- related to idea of "spurious wakeup"



# Question

---

The wait() call on a condition variable requires a parameter for a condition variable and a parameter for a lock.

- a) true
- b) false

True.

# Question

---

When a thread calls `wait()`, the lock it passes must be in which state?

- a) acquired
- b) free

acquired. The thread must be holding the lock when it calls `wait()`.

# Question

---

If your code is designed following the Anderson/Dahlin guidelines, then your code will never break if a 'spurious wakeup' occurs

- a) true
- b) false

True. A spurious wakeup is when a call to `wait()` returns even if a signal has not occurred.

# Question

---

In the Anderson/Dahlin style of designing shared objects, each object has one lock

- a) true
- b) false

True. That lock is used to ensure mutually exclusive access to shared variables.

# Semaphores

---

A semaphore is like a variable, but with three differences:

1. when you create a semaphore, you can initialize it to any value, but afterward there are only two operations you can perform
  - increment
  - decrement
2. when a thread decrements the semaphore, if the result is negative, the thread blocks itself and can't continue until another thread increments the semaphore
3. when a thread increments the semaphore, if there are threads waiting, one of them gets unblocked

Semaphores were invented by Edsger Dijkstra

# 2-thread rendezvous with semaphores

---

```
sem1 = Semaphore(0)  
sem2 = Semaphore(0)
```

thread A

```
1  a1  
2  sem1.signal()  
3  sem2.wait()  
4  a2
```

thread B

```
1  b1  
2  sem2.signal()  
3  sem1.wait()  
4  b2
```

# Mutual exclusion with semaphores

---

```
sem = Semaphore(1)
```

thread A

```
sem.wait()  
    // critical section  
    count = count + 1  
sem.signal()
```

thread B

```
sem.wait()  
    // critical section  
    count = count + 1  
sem.signal()
```

A “symmetric” solution

# Question

---

initialization:

```
sem = Semaphore(0)
```

thread A

```
sem.signal()  
print 'A'
```

thread B

```
sem.wait()  
print 'B'
```

What do we know about the output of this program:

- a) 'A' must be printed before 'B'
- b) 'B' must be printed before 'A'
- c) 'A' and 'B' are both printed, in either order

c. The `signal()` call does not mean that B starts running immediately.



# Make sure you understand:

---

concept of “atomic execution”

use of synchronization variables in Dahlin method

your solutions to concurrency homeworks

how to do address translation with paging

how caching affects memory access time

how to write awk scripts, esp. ones that use associative arrays

how to use find

how to write Make files