

I/O Devices

Glenn Bruns
CSUMB

Lecture Objectives

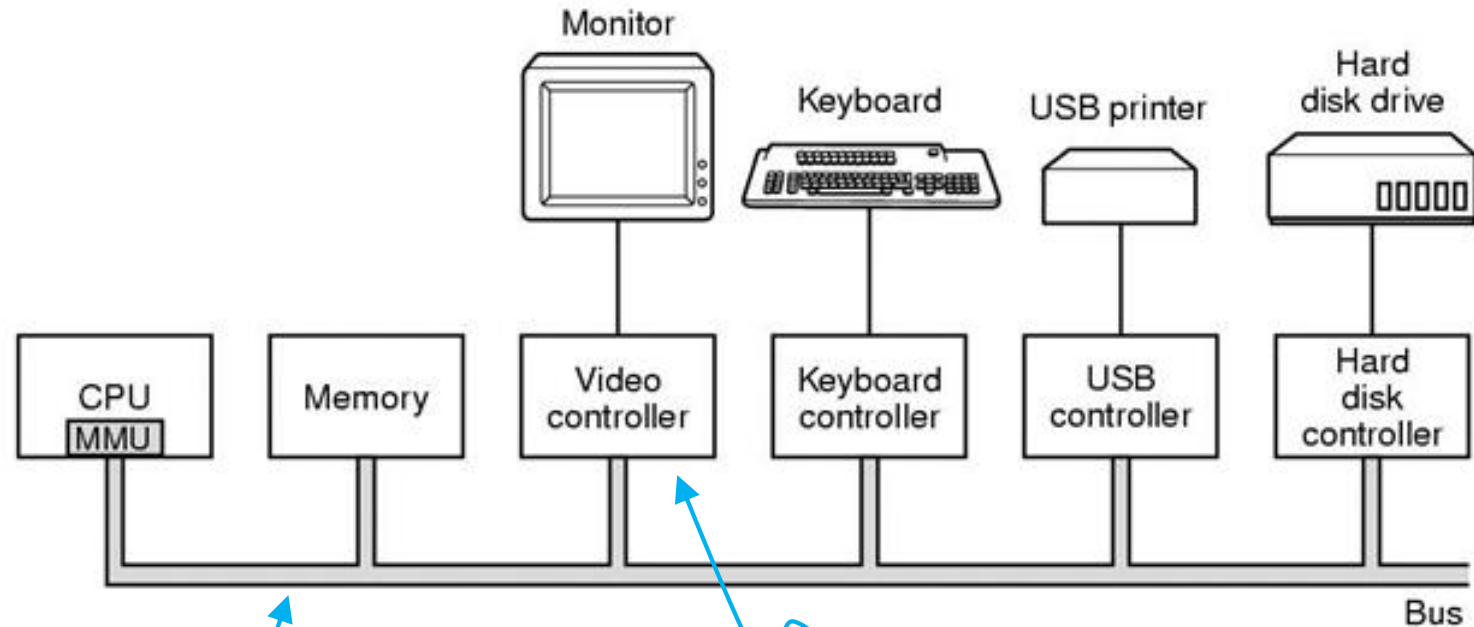
After this lecture, you should be able to:

- ❑ list the characteristics of block and character devices
- ❑ describe the hardware interfaces of I/O devices
- ❑ explain how the OS interacts with I/O devices

Motivation

- How does the OS actually talk to devices?
 - what CPU operations are involved?
 - how are devices addressed?
- How does the OS abstract away from the details of specific devices?
 - what is the software architecture?

Review: computer architecture

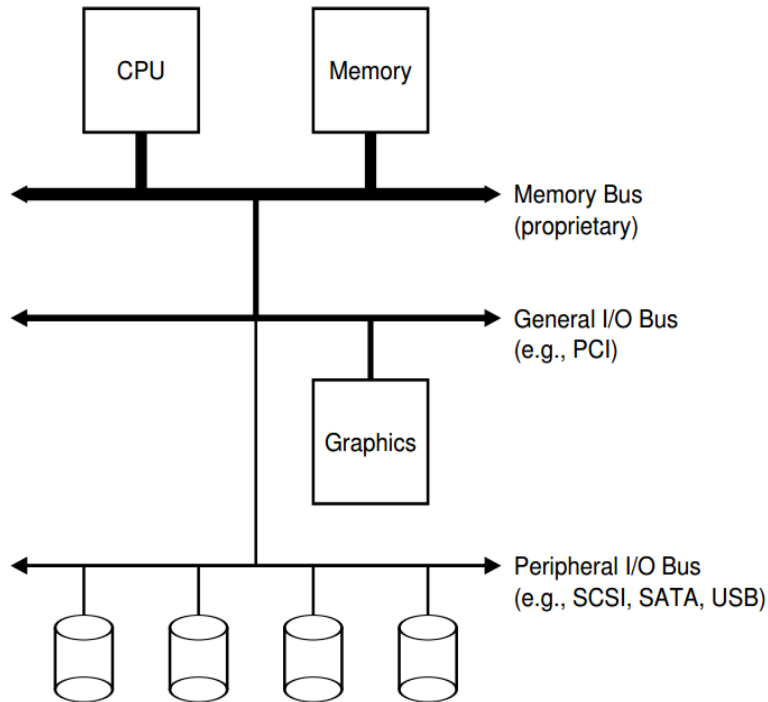


the bus carries bytes between components

a controller is hardware; it can be located in the computer or integrated into a device

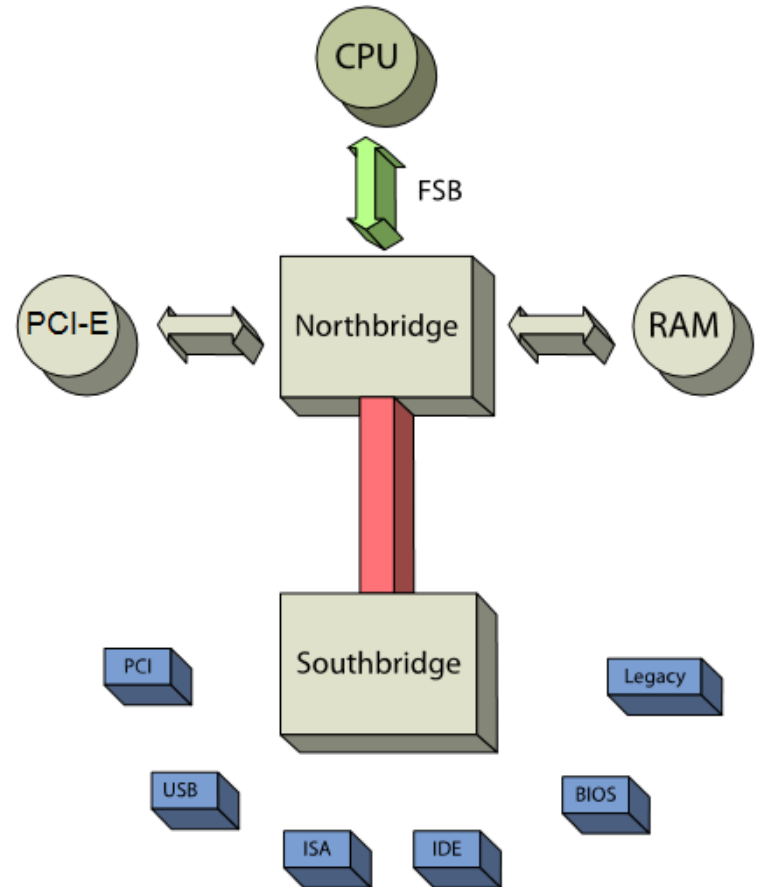
(Figure from Tanenbaum, *Modern Operating Systems*, 3rd edition)

Memory and I/O buses



(Source: Arpacci-Dusseau, *OSTEP*, chapter 36)

PC architecture



(Source: By Fred the Oyster, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=35241077>)

I/O Devices

□ Block devices

- example: hard drive
- random access of fixed-size “blocks” of data
- each block has an address
- block size typically 512 bytes to 32 KB
- transfers are in units of entire blocks

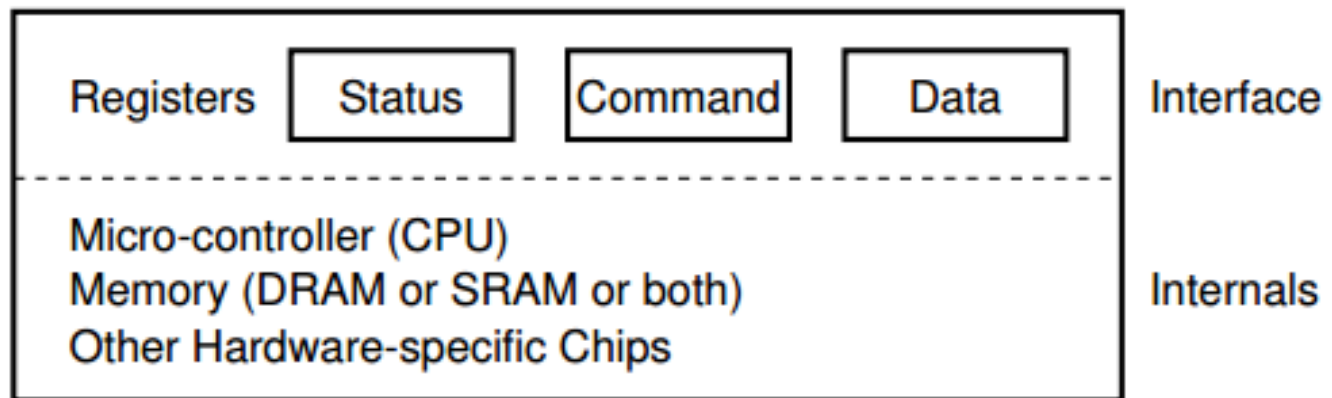
□ Character devices

- examples: printers, mice, keyboards
- processes a stream of bytes, one after another
- character devices are not addressable

I/O Device structure

What's in a device?


- ❑ the same things you might find in any computer
 - CPU, memory, special purpose chips
- ❑ a hardware interface
 - typically a set of **registers**
 - status register: read it to see device status
 - command register: write to it to tell device to do something
 - data register – used to pass data to/from device



How CPU interacts with device


```
while (status == busy)
;
write data to data register
write command to command register
while (status == busy)
;
```

executes
command



This is called “programmed I/O” – the CPU is running code to talk to the device.

Note use of polling.



Option 2: Interrupt-driven I/O

To avoid polling, I/O can be interrupt-driven:

- ❑ OS issues an I/O request
- ❑ OS puts calling process to sleep
- ❑ when I/O device is done, it raises a hardware interrupt
- ❑ interrupt handler finishes processing of request; wakes the waiting process





figure source: OSTEP

Problem with data movement



Option 3: Direct Memory Access (DMA)

DMA uses a special hardware device that transfers data between devices and memory without CPU help.

- ☐ CPU tells DMA about device, memory location, quantity to transfer,...
- ☐ DMA does the data transfer
- ☐ when done, DMA raises an interrupt

How does OS talk to devices?

Hardware interface to devices usually has registers.

How does CPU read/write those registers?


Approach 1:

- CPU has special instructions, such as 'in', 'out'
- the arguments specify a CPU register, and a "port" that names a device

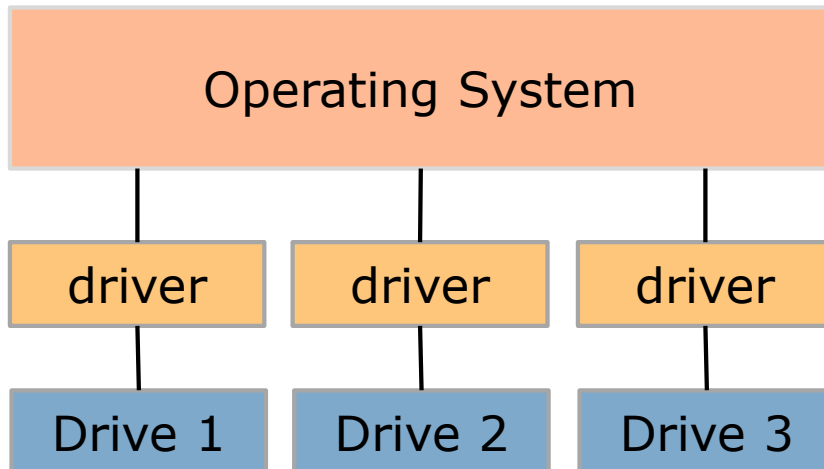
Approach 2:

- device registers are mapped to memory addresses ("memory-mapped I/O")

Insulating OS code from hardware



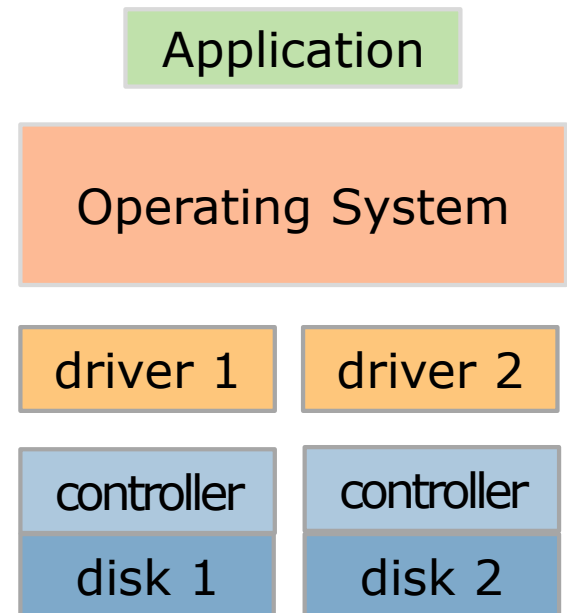
Bad: OS is aware of all types of driver interface



Good: OS is oblivious to differences between drives

Device drivers

- ❑ Each device type needs its own driver code
- ❑ There are lots of devices, so lots of drivers
- ❑ 70% of Linux is device driver code!
- ❑ Microsoft hates device drivers
 - many Windows crashes caused by bad device drivers
 - lots of money on software to automatically find problems with driver code



Hardware interface for IDE disk drive

Control Register:

Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means "enable interrupt"

Command Block Registers:

Address 0x1F0 = Data Port

Address 0x1F1 = Error

Address 0x1F2 = Sector Count

Address 0x1F3 = LBA low byte

Address 0x1F4 = LBA mid byte

Address 0x1F5 = LBA hi byte

Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive

Address 0x1F7 = Command/status

logical block address



Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

Error Register (Address 0x1F1): (check when Status ERROR==1)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF

BBK = Bad Block

UNC = Uncorrectable data error

MC = Media Changed

IDNF = ID mark Not Found

MCR = Media Change Requested

ABRT = Command aborted

T0NF = Track 0 Not Found

AMNF = Address Mark Not Found

Device driver for IDE disk drive

Control Register:

Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means

Command Block Registers:

Address 0x1F0 = Data Port

Address 0x1F1 = Error

Address 0x1F2 = Sector Count

Address 0x1F3 = LBA low byte

Address 0x1F4 = LBA mid byte

Address 0x1F5 = LBA hi byte

Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive

Address 0x1F7 = Command/status

Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

Error Register (Address 0x1F1): (check when Status ERROR)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF

BBK = Bad Block

UNC = Uncorrectable data error

MC = Media Changed

IDNF = ID mark Not Found

MCR = Media Change Requested

ABRT = Command aborted

T0NF = Track 0 Not Found

AMNF = Address Mark Not Found

Basic idea for a write:

- wait for drive to be ready (status register)
- write LBAs, drive number to command registers
- start the I/O (command register 0x1f7)
- data transfer (wait for READY and DRQ; write data to data port)
- handle interrupts
- error handling (check status register)

Summary

- ❑ The hardware interface of I/O devices normally consists of some registers (status, data, command,...)
- ❑ The OS reads/writes these registers via either:
 - special CPU I/O instructions (in/out)
 - memory-mapped I/O
- ❑ Device drivers insulate the OS from details of individual devices